

occupancy-dataset

June 10, 2023

1 ROOM OCCUPANCY ESTIMATION

Classification

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
[2]: df = pd.read_csv("C:/Users/Aditya Singh/Downloads/archive/Occupancy_Estimation.
↳csv")
df.head()
```

```
[2]:
```

	Date	Time	S1_Temp	S2_Temp	S3_Temp	S4_Temp	S1_Light	\
0	2017/12/22	10:49:41	24.94	24.75	24.56	25.38	121	
1	2017/12/22	10:50:12	24.94	24.75	24.56	25.44	121	
2	2017/12/22	10:50:42	25.00	24.75	24.50	25.44	121	
3	2017/12/22	10:51:13	25.00	24.75	24.56	25.44	121	
4	2017/12/22	10:51:44	25.00	24.75	24.56	25.44	121	

	S2_Light	S3_Light	S4_Light	S1_Sound	S2_Sound	S3_Sound	S4_Sound	\
0	34	53	40	0.08	0.19	0.06	0.06	
1	33	53	40	0.93	0.05	0.06	0.06	
2	34	53	40	0.43	0.11	0.08	0.06	
3	34	53	40	0.41	0.10	0.10	0.09	
4	34	54	40	0.18	0.06	0.06	0.06	

	S5_CO2	S5_CO2_Slope	S6_PIR	S7_PIR	Room_Occupancy_Count
0	390	0.769231	0	0	1
1	390	0.646154	0	0	1
2	390	0.519231	0	0	1
3	390	0.388462	0	0	1
4	390	0.253846	0	0	1

```
[3]: df.shape
```

```
[3]: (10129, 19)
```

```
[4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10129 entries, 0 to 10128
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                   10129 non-null  object
1   Time                   10129 non-null  object
2   S1_Temp                10129 non-null  float64
3   S2_Temp                10129 non-null  float64
4   S3_Temp                10129 non-null  float64
5   S4_Temp                10129 non-null  float64
6   S1_Light               10129 non-null  int64
7   S2_Light               10129 non-null  int64
8   S3_Light               10129 non-null  int64
9   S4_Light               10129 non-null  int64
10  S1_Sound               10129 non-null  float64
11  S2_Sound               10129 non-null  float64
12  S3_Sound               10129 non-null  float64
13  S4_Sound               10129 non-null  float64
14  S5_CO2                 10129 non-null  int64
15  S5_CO2_Slope           10129 non-null  float64
16  S6_PIR                 10129 non-null  int64
17  S7_PIR                 10129 non-null  int64
18  Room_Occupancy_Count  10129 non-null  int64
dtypes: float64(9), int64(8), object(2)
memory usage: 1.5+ MB

```

```

[16]: df['Date'] = pd.to_datetime(df['Date'])
df['Date'] = df['Date'].view('int64').astype('float64')
df['Time'] = pd.to_datetime(df['Time'], format='%H:%M:%S').dt.time
df['Time'] = df['Time'].apply(lambda x: x.hour * 3600 + x.minute * 60 + x.
↪second).astype('float64')

```

```

[20]: df.info()

```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10129 entries, 0 to 10128
Data columns (total 19 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                   10129 non-null  float64
1   Time                   10129 non-null  float64
2   S1_Temp                10129 non-null  float64
3   S2_Temp                10129 non-null  float64
4   S3_Temp                10129 non-null  float64
5   S4_Temp                10129 non-null  float64
6   S1_Light               10129 non-null  int64
7   S2_Light               10129 non-null  int64

```

```

8   S3_Light          10129 non-null  int64
9   S4_Light          10129 non-null  int64
10  S1_Sound          10129 non-null  float64
11  S2_Sound          10129 non-null  float64
12  S3_Sound          10129 non-null  float64
13  S4_Sound          10129 non-null  float64
14  S5_CO2            10129 non-null  int64
15  S5_CO2_Slope      10129 non-null  float64
16  S6_PIR            10129 non-null  int64
17  S7_PIR            10129 non-null  int64
18  Room_Occupancy_Count 10129 non-null  int64
dtypes: float64(11), int64(8)
memory usage: 1.5 MB

```

```
[21]: df.isnull().sum()
```

```

[21]: Date          0
      Time          0
      S1_Temp       0
      S2_Temp       0
      S3_Temp       0
      S4_Temp       0
      S1_Light      0
      S2_Light      0
      S3_Light      0
      S4_Light      0
      S1_Sound      0
      S2_Sound      0
      S3_Sound      0
      S4_Sound      0
      S5_CO2        0
      S5_CO2_Slope  0
      S6_PIR        0
      S7_PIR        0
      Room_Occupancy_Count 0
      dtype: int64

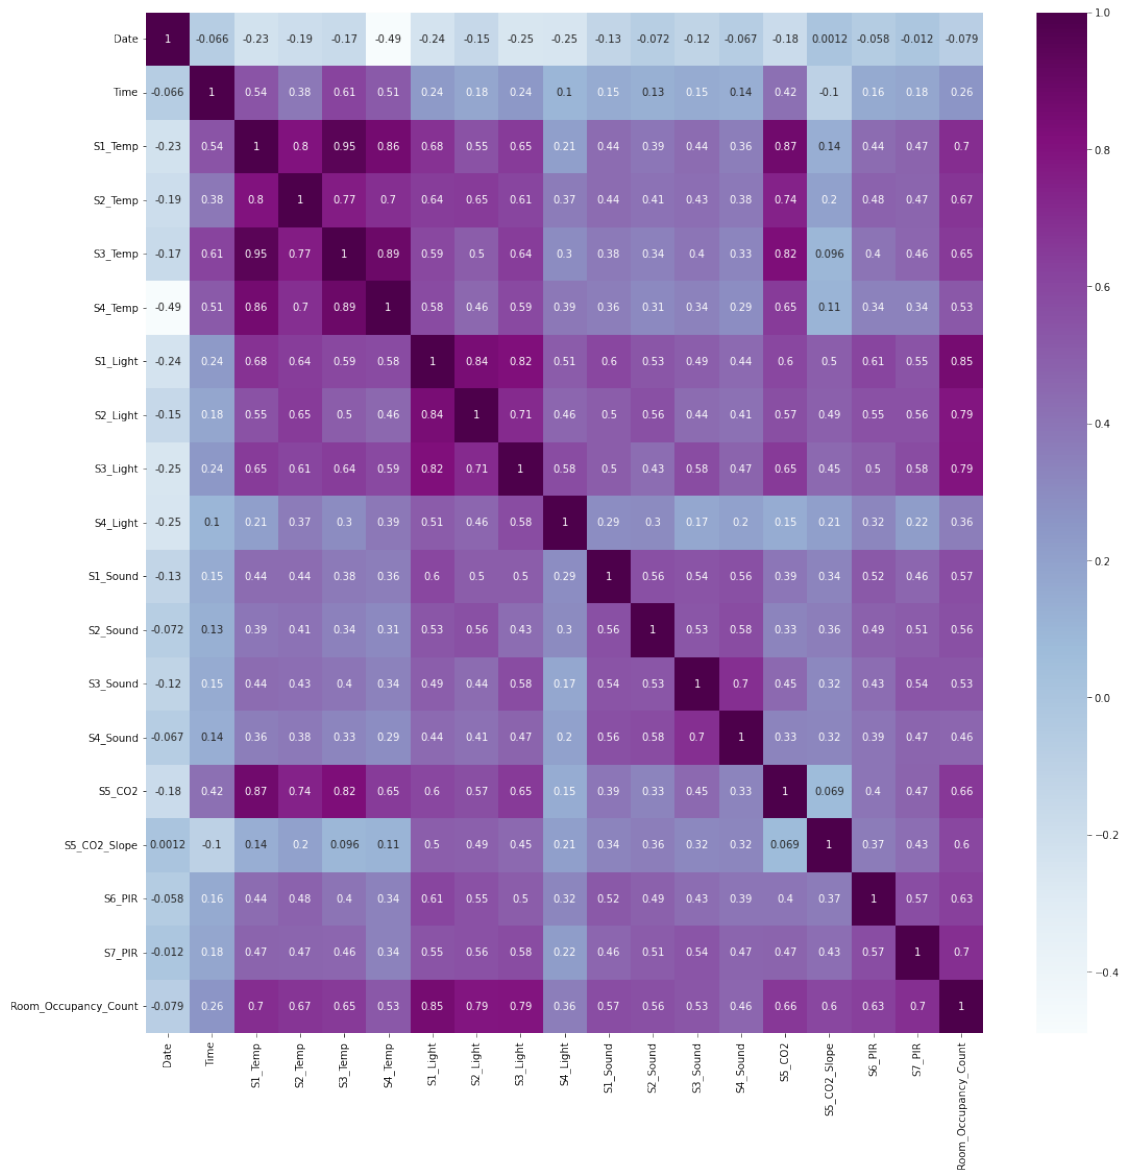
```

```

[22]: import seaborn as sns
      plt.figure(figsize=(18,18))
      sns.heatmap(df.corr(),cbar=True,annot=True,cmap="BuPu")

```


```
[22]: <AxesSubplot:>
```



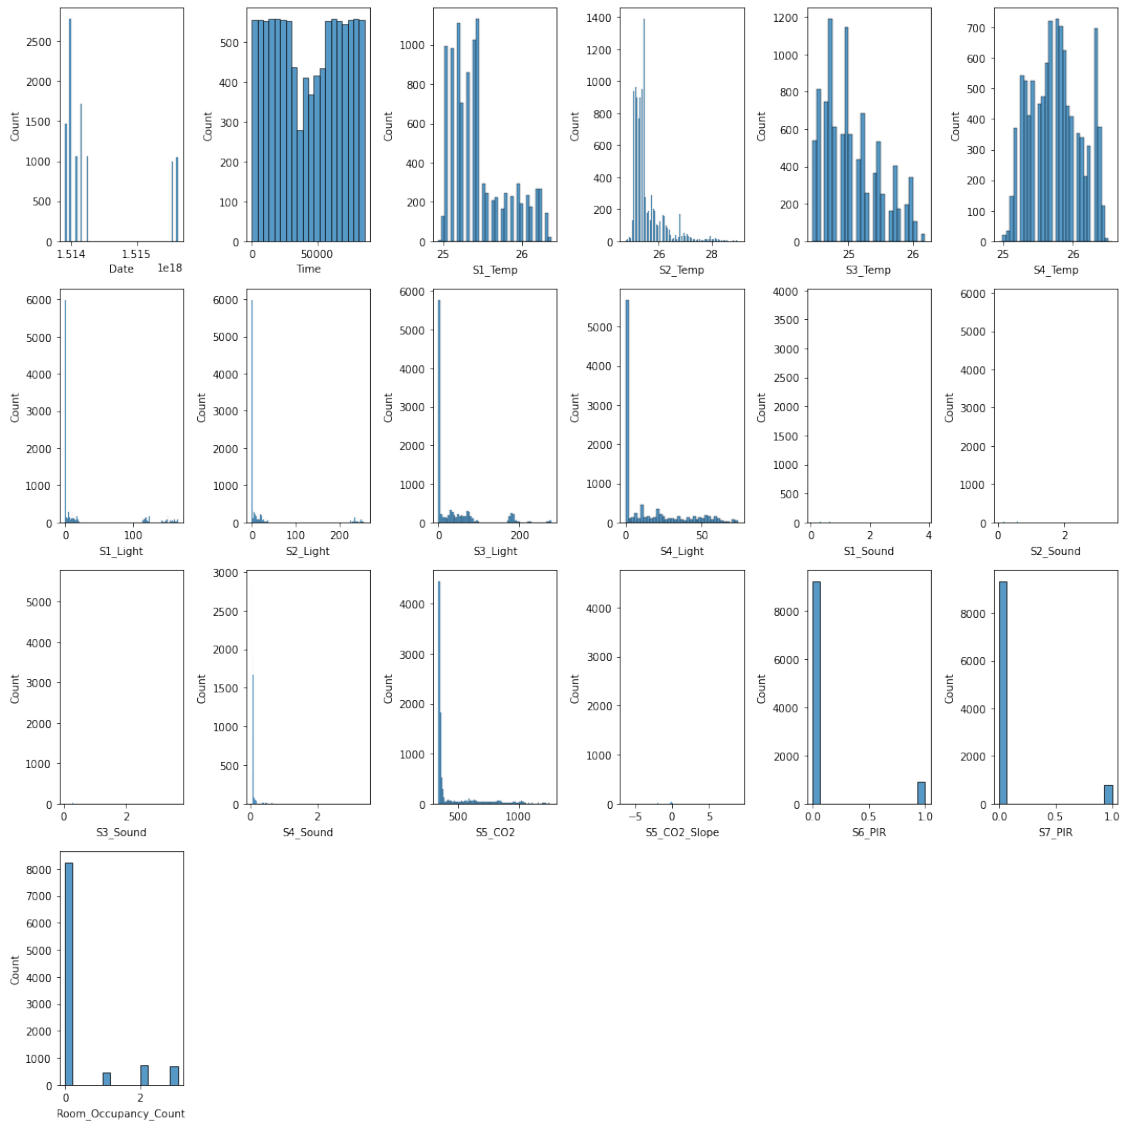
```
[23]: plt.figure(figsize=(18, 18))
sns.pairplot(df.corr())
```

```
[23]: <seaborn.axisgrid.PairGrid at 0x138b6c11550>
```

```
<Figure size 1296x1296 with 0 Axes>
```



```
[24]: count=1
plt.subplots(figsize=(15, 15))
for i in df.columns:
    plt.subplot(4,6,count)
    sns.histplot(df[i])
    count+=1
plt.tight_layout()
plt.show()
```



```
[25]: x = df.drop(['Room_Occupancy_Count'],axis=1).values
      y = df[['Room_Occupancy_Count']].values
```

```
[26]: from sklearn.model_selection import train_test_split
      x_train,x_test,y_train,y_test = train_test_split(x,y,test_size = 0.
      ↪3,random_state=0)
```

```
[27]: print(x_train.shape)
      print(y_train.shape)
      print(x_test.shape)
      print(y_test.shape)
```

(7090, 18)

```
(7090, 1)
(3039, 18)
(3039, 1)
```

```
[28]: from sklearn.linear_model import LogisticRegression
      from sklearn.svm import SVC
      from sklearn.neighbors import KNeighborsClassifier
      from sklearn.naive_bayes import GaussianNB
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.metrics import accuracy_score

      clf1 = LogisticRegression()
      clf2 = SVC()
      clf3 = KNeighborsClassifier()
      clf4 = GaussianNB()
      clf5 = DecisionTreeClassifier()

      clf = [clf1, clf2, clf3, clf4, clf5]
      clf_name = ['LR', 'SVC', 'KNN', 'GNB', 'DT']
      acc = {}

      for model, model_name in zip(clf, clf_name):
          model.fit(x_train, y_train)
          pred = model.predict(x_test)
          acc[model_name] = accuracy_score(y_test, pred) * 100

      print("ACCURACY SCORES")
      for i, j in acc.items():
          print(i, ':-', j, '%')

      print(acc.keys())
      print(acc.values())

      plt.figure()
      sns.barplot(x=list(acc.keys()), y=list(acc.values()))
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
```

```
    return f(*args, **kwargs)
C:\ProgramData\Anaconda3\lib\site-
packages\sklearn\linear_model\_logistic.py:763: ConvergenceWarning: lbfgs failed
to converge (status=2):
ABNORMAL_TERMINATION_IN_LNSRCH.
```

Increase the number of iterations (max_iter) or scale the data as shown in:
<https://scikit-learn.org/stable/modules/preprocessing.html>


Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    return f(*args, **kwargs)
C:\ProgramData\Anaconda3\lib\site-
packages\sklearn\neighbors\_classification.py:179: DataConversionWarning: A
column-vector y was passed when a 1d array was expected. Please change the shape
of y to (n_samples, ), for example using ravel().
    return self._fit(X, y)
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\validation.py:63:
DataConversionWarning: A column-vector y was passed when a 1d array was
expected. Please change the shape of y to (n_samples, ), for example using
ravel().
    return f(*args, **kwargs)

ACCURACY SCORES
LR :- 81.34254689042449 %
SVC :- 81.34254689042449 %
KNN :- 81.34254689042449 %
GNB :- 81.34254689042449 %
DT :- 99.53932214544258 %
dict_keys(['LR', 'SVC', 'KNN', 'GNB', 'DT'])
dict_values([81.34254689042449, 81.34254689042449, 81.34254689042449,
81.34254689042449, 99.53932214544258])
```

[28]: <AxesSubplot:>




```
[29]: base_pred = clf5.predict(x_test)
```

```
[30]: from sklearn.metrics import confusion_matrix , classification_report, \
      plot_confusion_matrix
      confusion_matrix(y_test,base_pred)
```

```
[30]: array([[2472,    0,    0,    0],
            [   0,  128,    3,    0],
            [   0,    2,  232,    5],
            [   2,    0,    2,  193]], dtype=int64)
```

```
[31]: plot_confusion_matrix(clf2,x_test,y_test)
```

```
[31]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x138badb4c40>
```




```
[32]: print(classification_report(y_test,base_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2472
1	0.98	0.98	0.98	131
2	0.98	0.97	0.97	239
3	0.97	0.98	0.98	197
accuracy			1.00	3039
macro avg	0.98	0.98	0.98	3039
weighted avg	1.00	1.00	1.00	3039

```
[33]: plt.figure(figsize = (7,7))
sns.heatmap(confusion_matrix(y_test,base_pred), annot = True, cbar = True, cmap=
↪ "BuPu", linewidths='1', linecolor = 'red')
```

```
[33]: <AxesSubplot:>
```



[]: